Shengwei Liu

New York, NY | 646-217-5499 | sl3597@cornell.edu
https://www.linkedin.com/in/shengwei127/

EDUCATION

Cornell Tech, Cornell University   New York, USA
MEng in Electrical and Computer Engineering, GPA: 4.019/4.0, Merit Scholarship   May 2026
Relevant Coursework: ASIC, Computer Architecture, Machine Learning, ML Hardware & Systems, Memory-Centric Computing
University College London (UCL)   London, UK
BEng in Electronic and Electrical Engineering, First-Class Honors   June 2025

TECHNICAL SKILLS

Programming: Python, C, C#, C++, CUDA, Embedded C, SystemVerilog, MATLAB
Frameworks & Libraries: TensorFlow, PyTorch, Keras, scikit-learn, SciPy, NumPy, Pandas, OpenAI Gym, TensorFlow Lite
EDA Tools: Synopsys (Design Compiler, PrimeTime), Cadence (Virtuoso, Innovus), Siemens EDA (QuestaSim, Calibre)
Operating Systems: macOS, Linux, Windows

EXPERIENCE

Cornell University, **Computer System Lab**, NY, USA   Dec. 2025 – Present
Research Assistant   Advisor: Prof. Mohamed Abdelfattah
- Investigated hardware–software co-designed LLM inference acceleration, focusing on low-bit quantization and numeric formats (FP4/MXFP4/NVFP4), decode-stage bottlenecks, and memory bandwidth efficiency.
- Analyzed architectural design choices for ASIC-style compute pipelines and PIM-based memory systems for low-bit inference workloads, identifying scalability limits under memory-bound conditions.

Tsinghua University, **HAS Lab**, Beijing, China   May 2025 – Sept. 2025
Research Assistant   Advisor: Prof. Shuwen Deng
- Implemented a YCSB-based benchmarking framework to systematically compare CXL-enabled memory systems against Redis cluster deployments, focusing on latency, throughput, and scalability under high-concurrency workloads.
- Analyzed performance implications of memory disaggregation and interconnect topologies in distributed KV stores, focusing on architectural trade-offs under varying concurrency levels.
- Evaluated Paxos/Raft variants optimized for CXL shared memory and RDMA, reducing replication latency by bypassing kernel networking stacks.
- Studied LSM-tree optimizations for CXL-based memory systems, analyzing effects of compaction strategy, index placement, and disaggregated memory on performance.
- This work led to a peer-reviewed publication accepted to IEEE ICDE 2026.

PROJECTS

GPT-2 Inference Kernel Optimization (C)   Nov. 2025 - Dec. 2025
- Implemented GPT-2 inference and performed microarchitectural profiling using Cachegrind, gprof, and gcov.
- Optimized the linear layer kernel using SIMD (AVX2/FMA), tiling, loop unrolling, memory flattening, and data-layout restructuring, achieving 3.95× speedup and 10× reduction in instruction count.
- Improved cache locality and memory bandwidth utilization by eliminating double-pointer indirection and applying data-reuse strategies.

Lightweight Pipelined Encryption Engine (SystemVerilog)   Oct. 2025 – Dec. 2025
- Designed a fully pipelined Simon32/64 lightweight block-cipher engine, exploring 1/2/4/8/16/32-stage microarchitectures to optimize throughput under strict PPA (power, performance, and area) constraints.
- Implemented complete RTL, performed functional verification across behavioral, post-synthesis, and post-layout netlists (QuestaSim), and achieved functional correctness over 1,000 randomized test vectors.
- Completed full ASIC flow including synthesis (Synopsys DC), automatic place-and-route with clock-tree synthesis (Cadence Innovus), and timing closure using post-route STA (PrimeTime).
- Generated final GDSII, resolved DRC/LVS issues (Calibre), and analyzed post-layout performance and power consumption.

Optimizing ML Kernel Operations (C)   Aug. 2025 – Oct. 2025
- Built core ML kernels (MatMul, Conv, ReLU) from scratch to understand computational bottlenecks.
- Applied advanced optimization techniques including tiling, blocking, and mixed-precision arithmetic to enhance performance.
- Explored multi-threading and parallel execution to improve CPU utilization.
- Integrated sparsity-aware methods to accelerate sparse matrix operations in ML inference.
- Combined all optimized kernels into an end-to-end neural network inference benchmark pipeline and evaluated performance.

Hand Motion Detection using IMU and PPG Sensors (Python, Embedded C)   Oct. 2024 – Apr. 2025
- Developed a real-time hand motion detection system integrating 6-axis Inertial Measurement Unit and 4-channel Photoplethysmography sensors, enabling Bluetooth-based multi-modal data streaming and edge-side preprocessing.
- Collected 1,800 samples across 6 hand motions, each with 15-time steps and 10 fused features, and utilized SciPy, Keras, and Scikit-learn for preprocessing and feature extraction, including Min-Max normalization, signal filtering, and stratified splitting.
- Trained a stacked LSTM network (25-15-10 units) with dropout and batch normalization via TensorFlow, achieving 94.05% test accuracy and 0.94 macro F1-score, outperforming SVM, KNN, and Random Forest baselines.
- Deployed quantized LSTM models on an Arduino Nano 33 IoT using the LiteRT framework and implemented a Python-based BLE inference pipeline with the Bleak package for real-time motion feedback to a PC without data loss.

PUBLICATION

Shengwei Liu, et al. "SHMemora: Protective Key–Value Store on Distributed Shared Memory". In: *Proceedings of the 42nd IEEE International Conference on Data Engineering* (**ICDE'26**), May 2026